

Files and File Systems

Major OS Themes

Virtualization

- Present physical resource as a more general, powerful, or easy-to-use form of itself
- Present illusion of multiple (or unlimited) resources where only one (or a few) really exist
- Examples: CPU, Memory, **Hard drives**

Concurrency

- Coordinate multiple activities to ensure correctness

Persistence


- Some data needs to survive crashes and power failures
-
- Need abstractions, mechanisms, policies for all

How to virtualize persistent storage?

- Recall OS Goals:
 - **Convenience** for the user and **efficient** use of machine
- Virtualization and abstraction helps with first goal
 - Files and directories abstract away the hard drive
- Efficiency:
 - File system controls when and how data is transferred to persistent storage

Storage Hierarchy

- processor registers, main memory, and auxiliary memory form a rudimentary memory hierarchy
- the hierarchy can be classified according to memory speed, cost, and volatility
- caches can be installed to hide performance differences when there is a large access-time gap between two levels



Outline

- Persistent storage of data
- Files
- Directories
- File Operations
- Inodes
- Traversing Paths

File Systems

- Provide long-term information storage
- Requirements:
 - Store very large amounts of information
 - Information must survive the termination of process using it
 - Multiple processes must be able to access info concurrently
- Two views of file systems:
 - User view – convenient logical organization of information
 - OS view – managing physical storage media, enforcing access restrictions

File (Management) Systems

- Implement an abstraction ([files](#)) for secondary storage
- Organize files logically ([directories](#))
- Permit sharing of data between processes, people, and machines
- Protect data from unwanted access (security)

File Operations

- Creation
 - Find space in file system, add entry to **directory**
 - map **file name** to **location and attributes**
- Writing
- Reading
 - Dominant abstraction is “**file as stream**”
- Repositioning within a file
- Deleting a file
- Truncation and appending
 - May erase the contents (or part of the contents) of a file while keeping attributes

Handling File Operations

- Must search the directory for the entry associated with the named file
- When the file is first used, store its attribute info in a system-wide open-file table
 - The index into the open-file table is used on subsequent operations, so no searching is required

Unix example (open, read, write are syscalls):

```
main() {  
    char onebyte;  
    int fd = open("sample.txt", "r");  
    read(fd, &onebyte, 1);  
    write(STDOUT, &onebyte, 1);  
    close(fd);  
}
```

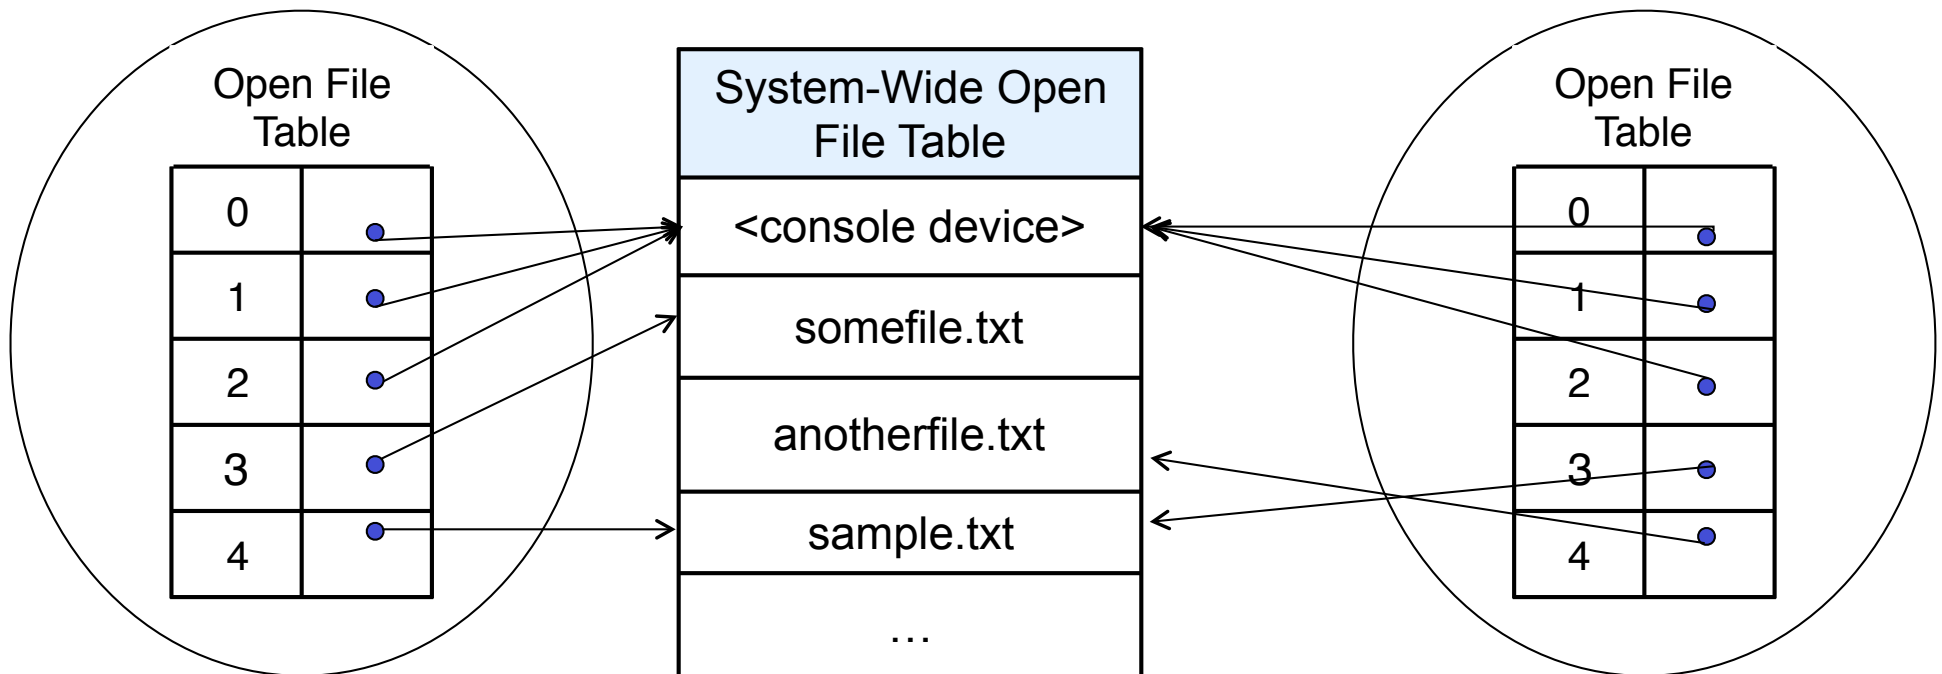
Open File Table
<console device>
...
sample.txt
...
...

File Sharing

- File sharing is incredibly important for getting work done
 - Basis for communication and synchronization
- Two key issues when sharing files
 - Semantics of concurrent access
 - What happens when one process reads while another writes?
 - What happens when two processes open a file for writing?
 - Protection

Shared Open Files

- There are two levels of internal tables
 1. A **per-process table** of all files that each process has open that contains the current file position for that process
 2. Each entry in the per-process table has a pointer to an entry in the **system-wide open-file table** for process independent info




File Access Methods

- General-purpose file systems support simple methods
 - Sequential access – read bytes one at a time, in order
 - Direct access – random access given block/byte number
- Database systems support more sophisticated methods
 - Record access - fixed or variable length
 - Indexed access
- What file access method(s) does Unix/Linux, Windows provide?
- Older systems provide more complicated methods
 - Modern systems typically only support simple access

Directories

- Directories provide logical structure to file systems
 - For users, they provide a means to organize files
 - For the file system, they provide a convenient naming interface
 - Allows the implementation to separate logical file organization from physical file placement
 - Stores information about files (owner, permission, etc.)
- Most file systems support multi-level directories
 - Naming hierarchies (/, /usr, /usr/local/, /home, ...)




Directory Structure

- A directory is a list of entries – names and associated metadata
 - Metadata is not the data itself, but information that describes properties of the data (size, protection, location, etc.)
- List is usually unordered (effectively random)
 - Entries usually sorted by program that reads directory
- Directories typically stored in files
 - Only need to manage one kind of secondary storage unit


Directory Implementations

- Single-level, two-level, or tree-structured
- Acyclic-graph directories: allows for shared directories
 - The same file or subdirectory may be in 2 different directories

Tree-structured:



Acyclic graph:



Directory Implementation

- Option 1: List
 - Simple list of file names and pointers to data blocks
 - Requires linear search to find entries
 - Easy to implement, slow to execute
 - And **directory operations are frequent!**
- Option 2: Hash Table
 - Create a list of file info structures
 - Hash file name to get a pointer to the file info structure in the list
 - Hash table takes space

File Links

Sharing can be implemented by creating a new directory entry called a **link**: a pointer to another file or subdirectory

- **Symbolic**, or **soft**, link
 - Directory entry refers to file that holds “true” path to the linked file
- **Hard** links
 - Second directory entry identical to the first

‘/’ directory

File Name	Start Block	Type
...
local	42	dir
usr	150	dir

‘usr’ directory (hard link) ‘usr’ directory (soft link)

File Name	Start Block	Type
...
local	42	dir
...

File Name	Start Block	Type
...
local	215	link
...

Issues with Links

- With links, a file may have multiple absolute path names
 - Traversing a file system should avoid traversing shared structures more than once
- Maintaining consistency is a problem
 - How do you update permissions in directory entry with a hard link?
- Deletion: When can the space allocated to a shared file be deallocated and reused?
 - Somewhat easier to handle with symbolic links
 - Deletion of a link is OK; deletion of the file entry itself deallocates space and leaves the link pointers dangling
 - Keep a reference count for hard links
- Sharing: How can you tell when two processes are sharing the same file?

Next Up

- Implementing File Systems
 - Disk layout
 - File metadata
 - Directory data


File System Implementation

- How do file systems use the disk to store files?
 - File systems define a **block size** (e.g., 4KB)
 - Disk space is allocated in granularity of blocks
- A “**Master Block**” determines location of root directory (aka *partition control block, superblock*)
 - Always at a well-known disk location
 - Often replicated across disk for reliability
- A **free map** determines which blocks are free
 - Usually a bitmap, one bit per block on the disk
 - Also stored on disk, cached in memory for performance
- Remaining disk blocks used to store files (and dirs)
 - There are many ways to do this

Disk Layout Strategies


- Files often span multiple disk blocks (think “processes often span multiple pages ...”)
- How do you find all of the blocks for a file?
 1. **Contiguous allocation**
 - Fast, simplifies directory access and allows indexing
 - Inflexible, causes external fragmentation, requires compaction
 2. **Linked**, or chained, structure
 - Each block points to the next, directory points to the first
 - Good for sequential (streaming) access, bad for all others
 3. **Indexed** structure (kind of like address translation)
 - An “index block” contains pointers to many other blocks
 - Handles random access better, still good for sequential
 - May require multiple, linked index blocks

Contiguous Allocation



File Name	Start Blk	Length
File A	2	3
File B	9	5
File C	18	8
File D	27	2

Linked Allocation



File Name	Start Blk	Last Blk
...
File B	1	22
...

Unix Inodes

- Unix **inodes** implement an indexed structure for files
- All file metadata is stored in an inode
 - Unix directory entries map file names to inodes
- Each inode contains 15 block pointers
 - First 12 are direct block pointers
 - Disk addresses of first 12 data blocks in file
 - The 13th is a single indirect block pointer
 - Address of block containing addresses of data blocks
 - Then the 14th is a double indirect block pointer
 - Address of block containing addresses of single indirect blocks
 - Finally, the 15th is a triple indirect block pointer

Example UNIX Inode

- Ext2 Linux file system Inodes are 72 bytes

